

Prediction of Default Risk at LoanTap using Logistic Regression

In the rapidly evolving digital lending landscape, Non-Banking Financial Companies (NBFCs) like LoanTap leverage technology to provide quick and accessible credit solutions on consumer friendly terms to salaried professionals and businessmen. A critical challenge for any lending institution is accurately assessing the creditworthiness of applicants to mitigate the inherent risk of loan defaults.

This case study explores the application of Logistic Regression, a foundational machine learning algorithm, to build a robust predictive model for loan default. By analyzing various applicant attributes and historical loan performance data, this project aims to identify key risk factors and predict the probability of a prospective borrower defaulting. The insights derived from this model will enable LoanTap to make more informed, data-driven lending decisions. This case study will focus on the underwriting process behind Personal Loans only.

Importing Libraries

```
In [259]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import datetime
import time
from sklearn.impute import KNNImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, roc_auc_score, precision_recall_curve, auc


warnings.simplefilter(action='ignore', category=Warning)
# setting the Custom Palette to have color coding consistent for all charts
custom_palette = ["#7eb0d5", "#fd7f6f", "#b2e061", "#bd7ebe", "#ffb55a"]
sns.set_palette(custom_palette)
```

Exploratory Data Analysis

```
In [260]: df=pd.read_csv("logistic_regression.csv")
pd.set_option('display.max_columns', None)
df.head()
```

Out[260]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_o
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	M
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	M



```
In [261]: # Look at the datatypes of the columns
print('#####')
print(f'Shape of the dataset is:{df.shape}')
print('#####')
print(df.info())
print('#####')
print(f'Number of nan/null values in each column: \n{df.isna().sum()}')
print('#####')
print(f'Number of duplicate entries: \n{df.duplicated().value_counts()}')
```

```
#####
Shape of the dataset is:(396030, 27)
#####
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null  float64
1   term                                396030 non-null  object
2   int_rate                            396030 non-null  float64
3   installment                         396030 non-null  float64
4   grade                               396030 non-null  object
5   sub_grade                          396030 non-null  object
6   emp_title                          373103 non-null  object
7   emp_length                         377729 non-null  object
8   home_ownership                     396030 non-null  object
9   annual_inc                        396030 non-null  float64
10  verification_status                396030 non-null  object
11  issue_d                           396030 non-null  object
12  loan_status                       396030 non-null  object
13  purpose                           396030 non-null  object
14  title                             394275 non-null  object
15  dti                               396030 non-null  float64
16  earliest_cr_line                  396030 non-null  object
17  open_acc                         396030 non-null  float64
18  pub_rec                         396030 non-null  float64
19  revol_bal                       396030 non-null  float64
20  revol_util                      395754 non-null  float64
21  total_acc                      396030 non-null  float64
22  initial_list_status              396030 non-null  object
23  application_type                 396030 non-null  object
24  mort_acc                       358235 non-null  float64
25  pub_rec_bankruptcies            395495 non-null  float64
26  address                         396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
None
#####
Number of nan/null values in each column:
loan_amnt                0
term                    0
int_rate                0
installment             0
grade                   0
sub_grade               0
emp_title               22927
emp_length              18301
home_ownership          0
annual_inc              0
verification_status     0
issue_d                 0
loan_status             0
purpose                 0
title                   1755
dti                     0
earliest_cr_line        0
open_acc                0
pub_rec                 0
revol_bal               0
revol_util              276
```

```
total_acc          0
initial_list_status 0
application_type    0
mort_acc           37795
pub_rec_bankruptcies 535
address            0
dtype: int64
#####
Number of duplicate entries:
False    396030
dtype: int64
```

```
In [262]: df.describe(include='all').T
```

Out[262]:

	count	unique	top	freq	mean	std
loan_amnt	396030.0	NaN	NaN	NaN	14113.888089	8357.441341
term	396030	2	36 months	302005	NaN	NaN
int_rate	396030.0	NaN	NaN	NaN	13.6394	4.472157
installment	396030.0	NaN	NaN	NaN	431.849698	250.72779
grade	396030	7	B	116018	NaN	NaN
sub_grade	396030	35	B3	26655	NaN	NaN
emp_title	373103	173105	Teacher	4389	NaN	NaN
emp_length	377729	11	10+ years	126041	NaN	NaN
home_ownership	396030	6	MORTGAGE	198348	NaN	NaN
annual_inc	396030.0	NaN	NaN	NaN	74203.175798	61637.621158
verification_status	396030	3	Verified	139563	NaN	NaN
issue_d	396030	115	Oct-2014	14846	NaN	NaN
loan_status	396030	2	Fully Paid	318357	NaN	NaN
purpose	396030	14	debt_consolidation	234507	NaN	NaN
title	394275	48817	Debt consolidation	152472	NaN	NaN
dti	396030.0	NaN	NaN	NaN	17.379514	18.019092
earliest_cr_line	396030	684	Oct-2000	3017	NaN	NaN
open_acc	396030.0	NaN	NaN	NaN	11.311153	5.137649
pub_rec	396030.0	NaN	NaN	NaN	0.178191	0.530671
revol_bal	396030.0	NaN	NaN	NaN	15844.539853	20591.836109
revol_util	395754.0	NaN	NaN	NaN	53.791749	24.452193
total_acc	396030.0	NaN	NaN	NaN	25.414744	11.886991
initial_list_status	396030	2	f	238066	NaN	NaN
application_type	396030	3	INDIVIDUAL	395319	NaN	NaN
mort_acc	358235.0	NaN	NaN	NaN	1.813991	2.14793
pub_rec_bankruptcies	395495.0	NaN	NaN	NaN	0.121648	0.356174
address	396030	393700	USCGC Smith\r\nFPO AE 70466	8	NaN	NaN



Understanding data distribution across categorical columns

```
In [263]: # Non-numeric columns
obj_cols = df.select_dtypes(include='object').columns
obj_cols=[item for item in obj_cols if item not in ['earliest_cr_line','issue_d']]

for col in obj_cols:
    print()
    print(f'Total Unique Values in {col} column are :- {df[col].nunique()}')
    percent=df[col].value_counts(normalize=True)*100
    percent=percent[percent>0.1]
    print(f'Value counts in {col} column are :-\n {percent}')
    print()
    print('-'*120)
```

Total Unique Values in term column are :- 2

Value counts in term column are :-

36 months	76.258112
-----------	-----------

60 months	23.741888
-----------	-----------

Name: term, dtype: float64

Total Unique Values in grade column are :- 7

Value counts in grade column are :-

B	29.295255
---	-----------

C	26.762366
---	-----------

A	16.207611
---	-----------

D	16.040199
---	-----------

E	7.950913
---	----------

F	2.972502
---	----------

G	0.771154
---	----------

Name: grade, dtype: float64

EDA Insights

- The dataset consists of 396030 records with 27 features, out of which one of them is a target variable.
- Loan_Status, which is the target, is a categorical variable.
- Out of all 26 features, 12 are numerical columns and 14 are object columns.
- There are no duplicates.
- There are 22927 missing values in emp_title, 18301 in emp_length, 1756 in title, 276 in revol_util, 37795 in mort_acc and 535 in pub_rec_bankruptcies.
- Based on their data types and number of unique values, the columns term, grade, sub_grade, emp_length, home_ownership, emp_title, title, verification_status, purpose, initial_list_status, pub_rec, and application_type, are categorical variables
- issue_d and earliest_cr_line can be converted to datetime and further month and year can be extracted.
- The term and emp_length column contains spaces and alphanumeric characters.
- Replace '36 months' with '36' and '60 months' with '60' in term column.
- Replace '<1 year' with 0 and '10 + years' with 10 in emp_length column.
- emp_title and title have lot of unique values and can be dropped as it will not have an impact on the loan approval.
- Extract zip code from address and address column can be dropped.
- As we already have grade column, we can drop sub-grade column

Transformation of columns:

```
In [264]: #convert to datetime
df['issue_year']=pd.to_datetime(df['issue_d']).dt.year.astype('object')
df['issue_month']=pd.to_datetime(df['issue_d']).dt.month.astype('object')
df['earliest_cr_line']=pd.to_datetime(df['earliest_cr_line']).dt.year.astype('object')

# extract zip code from address columns
df['zip_code']=df['address'].apply(lambda x:x[-5:])
```

Visual Analysis

Univariate Analysis











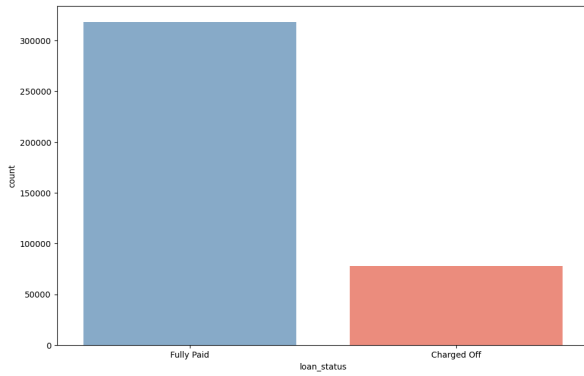
i.Categorical Variables



```
In [265]: cat_cols=['loan_status', 'term', 'grade', 'sub_grade',
                    'home_ownership', 'verification_status',
                    'emp_length', 'initial_list_status', 'application_type', 'zip_code', 'issue_d']
```



```
In [266]: cat_cols1=['loan_status','term', 'grade', 'sub_grade',
                    'home_ownership', 'verification_status',
                    'emp_length','initial_list_status','application_type','zip_code','issu
plt.figure(figsize=(20,50))
i=1
for col in cat_cols1:
    ax=plt.subplot(8,2,i)
    sns.countplot(df[col],ax=ax)
    i+=1

plt.tight_layout()
plt.show()
```








```
In [267]: cat_cols2=['purpose','earliest_cr_line']           #with lot of unique values

plt.figure(figsize=(15,20))
i=1
for col in cat_cols2:
    ax=plt.subplot(2,1,i)
    sns.countplot(df[col],ax=ax)
    plt.xticks(rotation=45)
    i+=1

plt.tight_layout()
plt.show()
```














- Majority of the loans are for 3 years.
- Most of the loans have grade B followed by Grade C.
- Most of the loans have sub-grade B3 followed by B4.
- Most of the loans have been taken by borrowers whose house is mortgaged, followed by rent.
- Most of the loans have been taken for debt consolidation followed by credit card.
- Most of the loans have initial list status as 'f'.
- Majority of the loans have been taken by individuals.
- Majority of the loans are verified.
- Majority of the loans have been taken by borrowers who have employment length of more than 10 years.

- Most of the loans are issued in October month, followed by July.

ii. Continuous Variables

```
In [268]: num_cols=df.select_dtypes(include='float64').columns
plt.figure(figsize=(20,50))
i=1
for col in num_cols:
    ax=plt.subplot(8,2,i)
    sns.histplot(df[col],ax=ax,kde=True,bins=30,color=custom_palette[3])
    i+=1
plt.tight_layout()
plt.show()
```
















Bivariate Analysis

i.Categorical Variables

```
In [280]: fig,axis= plt.subplots(5,2,figsize=(25,40))
sns.countplot(data=df,x='term',hue='loan_status',ax=axis[0,0])
sns.countplot(data=df,x='grade',hue='loan_status',ax=axis[0,1])
sns.countplot(data=df,x='home_ownership',hue='loan_status',ax=axis[1,0])
sns.countplot(data=df,x='application_type',hue='loan_status',ax=axis[1,1])
sns.countplot(data=df,x="initial_list_status",hue='loan_status',ax=axis[2][0])
sns.countplot(data=df,x="verification_status",hue='loan_status',ax=axis[2][1])
sns.countplot(data=df,x="emp_length",hue='loan_status',ax=axis[3][0])
sns.countplot(data=df,x="issue_month",hue='loan_status',ax=axis[3][1])
sns.countplot(data=df,x="zip_code",hue='loan_status',ax=axis[4][0])
sns.countplot(data=df,x="purpose",hue='loan_status',ax=axis[4][1])
plt.xticks(rotation=45)
plt.show()
```
















```
In [270]: plt.figure(figsize=(20,40))
i=1
for col in cat_cols:
    if col=='loan_status':
        continue
    ax=plt.subplot(8,2,i)

    data1=(pd.crosstab(df[col],df['loan_status'],normalize='index')*100)
    data1=data1.round(2)

    data1.plot(kind='bar', stacked=True,ax=ax)
    plt.xlabel(f'{col}')
    plt.ylabel('% Applicants')
    plt.title(f'% Defaulters by {col}')
    plt.legend(['Charged Off','Fully Paid'])
    i += 1

plt.tight_layout()
plt.show()
```




















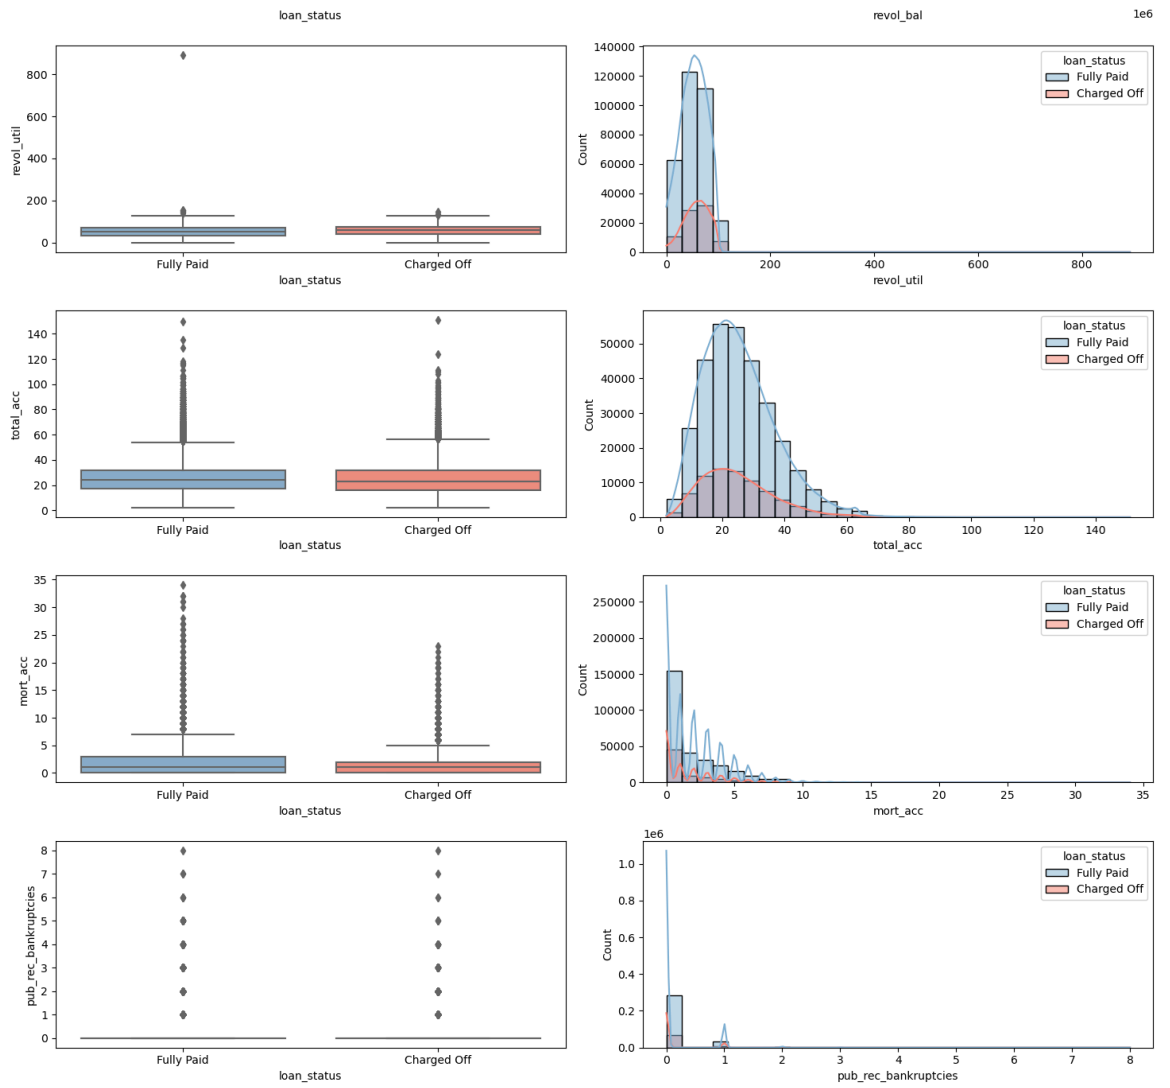
ii. Continuous Variables

```
In [281]: fig, ax = plt.subplots(12,2,figsize=(15,40))
          i=0
          for col in num_cols:

              sns.boxplot(data=df,x="loan_status",y=col,ax=ax[i,0])
              sns.histplot(data=df,x=col,hue='loan_status',ax=ax[i,1],kde=True, fill=True,
                           i += 1


          plt.tight_layout()
          plt.show()
```




Multivariate Analysis


```
In [272]: plt.figure(figsize=(7,5))
sns.scatterplot(data=df,x='loan_amnt',y='annual_inc',hue='loan_status')
plt.title('Annual Income vs Loan Amount',fontsize=10)
plt.xlabel('Annual Income',fontsize=8)
plt.ylabel('Loan Amount',fontsize=8)
plt.show()
```



```
In [273]: plt.figure(figsize=(7,5))
sns.scatterplot(data=df,x='loan_amnt',y='dti',hue='loan_status')
plt.title('DTI vs Loan Amount',fontsize=10)
plt.xlabel('Debt to Income',fontsize=8)
plt.ylabel('Loan Amount',fontsize=8)
plt.show()
```



```
In [274]: plt.figure(figsize=(7,5))
sns.scatterplot(data=df,x='loan_amnt',y='revol_bal',hue='loan_status')
plt.title('Revolving Balance vs Loan Amount',fontsize=10)
plt.xlabel('Revolving Balance',fontsize=8)
plt.ylabel('Loan Amount',fontsize=8)
plt.show()
```




```
In [275]: plt.figure(figsize=(25,17))
sns.heatmap(df[num_cols].corr(), cmap="GnBu",annot=True)
plt.show()
```



There seems to be a high correlation between loan amount and installment so we can drop one of them for our model building later.

Handling Missing Values

There are 22927 missing values in emp_title, 18301 in emp_length, 1756 in title, 276 in revol_util, 37795 in mort_acc and 535 in pub_rec_bankruptcies.

```
In [277]: df['emp_length']=df['emp_length'].replace({'< 1 year':'0', '10+ years':'10'})
df['emp_length']=df['emp_length'].str.replace(r'\D', '', regex=True)
```

```
In [278]: df['emp_title'].fillna('Unknown',inplace=True)
df['title'].fillna('Unknown',inplace=True)
```

```
In [282]: imputer=KNNImputer()
df['emp_length']=imputer.fit_transform(df[['emp_length']])
```

```
In [283]: mort_acc_mode=df.groupby('total_acc')['mort_acc'].agg(lambda x: pd.Series.mode
```

```
In [284]: def fill_mort(total_acc,mort_acc):  
          if np.isnan(mort_acc):  
              return mort_acc_mode[total_acc].round()  
          else:  
              return mort_acc
```

```
In [285]: df['mort_acc']=df.apply(lambda x: fill_mort(x['total_acc'],x['mort_acc']),axis=1)
```

```
In [286]: df.dropna(subset=['revol_util','pub_rec_bankruptcies'],inplace=True)
```










```
In [287]: df.isna().sum()
```

```
Out[287]: loan_amnt      0  
term      0  
int_rate  0  
installment  0  
grade     0  
sub_grade 0  
emp_title  0  
emp_length 0  
home_ownership  0  
annual_inc  0  
verification_status  0  
issue_d     0  
loan_status 0  
purpose     0  
title       0  
dti         0  
earliest_cr_line  0  
open_acc    0  
pub_rec     0  
revol_bal   0  
revol_util  0  
total_acc   0  
initial_list_status  0  
application_type  0  
mort_acc    0  
pub_rec_bankruptcies  0  
address     0  
issue_year  0  
issue_month 0  
zip_code    0  
dtype: int64
```

Outlier Treatment

We will not consider mort_acc, pub_rec & pub_rec_bankruptcies for outlier removal as we are going to take care of these columns later.










```
In [288]: outlier_cols=[item for item in num_cols if item not in ['mort_acc', 'pub_rec', '
plt.figure(figsize=(14,25))
i=1
for col in outlier_cols:
    ax=plt.subplot(5,2,i)
    sns.boxplot(x=df[col],ax=ax)
    i+=1
plt.show()
```

```
In [290]: #Using IQR method
# for col in outlier_cols:
#     Q1=df[col].quantile(0.25)
#     Q3=df[col].quantile(0.75)
#     IQR=Q3-Q1
#     Lower=Q1-1.5*IQR
#     upper=Q3+1.5*IQR
#     df=df[(df[col]>=Lower)&(df[col]<=upper)]
```

```
In [291]: # Using Z-Scores
for col in outlier_cols:
    mean=df[col].mean()
    stdev=df[col].std()
    lower=mean-3*stdev
    upper=mean+3*stdev
    df=df[(df[col]>=lower)&(df[col]<=upper)]
```

```
In [292]: plt.figure(figsize=(14,25))
i=1
for col in outlier_cols:
    ax=plt.subplot(6,2,i)
    sns.boxplot(x=df[col],ax=ax)
    i+=1
plt.show()
```



Data Preparation for Modelling

```
In [293]: X = df.drop(columns=['loan_status'])
y = df['loan_status']
y = y.replace({'Fully Paid': 0, 'Charged Off': 1}).astype(int)
```

Feature Engineering

Three columns have been modified as below:

```
In [294]: X['pub_rec'] = X['pub_rec'].apply(lambda x: 0 if x == 0 else 1)
X['mort_acc'] = X['mort_acc'].apply(lambda x: 0 if x == 0 else 1)
X['pub_rec_bankruptcies'] = X['pub_rec_bankruptcies'].apply(lambda x: 0 if x == 0 else 1)
```

```
In [296]: #transform term column
X['term'].replace({' 36 months': 36, ' 60 months': 60}, inplace=True)

#transform home_ownership column
X['home_ownership'].replace({'ANY': 'OTHER', 'NONE': 'OTHER'}, inplace=True)

#transform verification_status column
X['initial_list_status'].replace({'w': 0, 'f': 1}, inplace=True)

#drop these columns for model building
X.drop(columns=['issue_d', 'issue_year', 'earliest_cr_line', 'emp_title', 'title', '...
```

1. Encoding of Categorical Columns

```
In [297]: cat_columns = X.select_dtypes(include=['object']).columns
```

```
In [298]: encoder = OneHotEncoder(sparse=False, handle_unknown='ignore')
one_hot_encoded = encoder.fit_transform(X[cat_columns])
one_hot_df = pd.DataFrame(one_hot_encoded, columns=encoder.get_feature_names_out(X[cat_columns].columns))
X = pd.concat([X, one_hot_df], axis=1)
X.drop(cat_columns, axis=1, inplace=True)
```

```
In [299]: print(X.shape) # Should be (same number of rows, number of features)
print(y.shape) #
```

```
(298581, 67)
(298581,)
```

2. Splitting into Train and Test Data

```
In [300]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

3. Standardization


```
In [301]: scaler=StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train),columns = X_train.columns)
X_test = pd.DataFrame(scaler.transform(X_test),columns = X_test.columns)
```

4. SMOTE for imbalanced data

```
In [302]: y_train.value_counts(normalize=True)*100
```

```
Out[302]: 0    81.141151
          1    18.858849
          Name: loan_status, dtype: float64
```

```
In [303]: sm=SMOTE(random_state=42)
X_train,y_train=sm.fit_resample(X_train,y_train)
```

```
In [304]: y_train.value_counts(normalize=True)*100
```

```
Out[304]: 0    50.0
          1    50.0
          Name: loan_status, dtype: float64
```

Logistic Regression

```
In [305]: model = LogisticRegression()
model.fit(X_train,y_train)
```

```
Out[305]: LogisticRegression()
```

Model Performance Evaluation

```
In [306]: y_train_pred=pd.DataFrame(model.predict(X_train))
y_pred=pd.DataFrame(model.predict(X_test))
```

```
In [307]: report_train=classification_report(y_train,y_train_pred)
report_test=classification_report(y_test,y_pred)
print("Train Reports are: \n"+report_train)
print("Test Reports are: \n"+report_test)
```


Train Reports are:

	precision	recall	f1-score	support
0	0.81	0.79	0.80	193817
1	0.79	0.82	0.81	193817
accuracy			0.80	387634
macro avg	0.80	0.80	0.80	387634
weighted avg	0.80	0.80	0.80	387634


Test Reports are:

	precision	recall	f1-score	support
0	0.95	0.79	0.86	48491
1	0.47	0.81	0.60	11226
accuracy			0.79	59717
macro avg	0.71	0.80	0.73	59717
weighted avg	0.86	0.79	0.81	59717

```
In [308]: # Confusion Matrix
cm_train = confusion_matrix(y_train, y_train_pred)
disp = ConfusionMatrixDisplay(cm_train)
disp.plot()
plt.title('Confusion Matrix for Training Data')
plt.show()
```




```
In [309]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix for Testing Data')
plt.show()
```



Model Interpretability

```
In [310]: coefficients=model.coef_[0]
feature_names=X_train.columns
feature_importance_df=pd.DataFrame({'feature':feature_names,'importance':coeff
feature_importance_df=feature_importance_df.sort_values(by='importance',ascend
plt.figure(figsize=(8,8))
sns.barplot(y = feature_importance_df['feature'],
            x = feature_importance_df['importance'],color='orange')
plt.title("Feature Importance for Model")
plt.yticks(fontsize=8)
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```




ROC Curve & AUC

```
In [311]: # Make predictions on the test set
y_pred_proba = model.predict_proba(X_test)[: ,1]

# Compute ROC curve and ROC-AUC score
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' %
plt.plot([0, 1], [0, 1], color='blue', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```




Precision Recall Curve

```
In [312]: precision, recall, thr = precision_recall_curve(y_test, y_pred_proba)
```


```
In [313]: pr_auc = auc(recall, precision)

# Plot the precision-recall curve
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve (PR AUC = {:.2f})'.format(pr_auc))
plt.grid(True)
plt.show()
```



Plotting Precison & Recall at different Thresholds


```
In [314]: plt.figure()
plt.plot(thr,precision[0:len(thr)],label='precision',color='blue')
plt.plot(thr,recall[0:len(thr)],label='recall',color='orange')
intersection_thr = thr[np.argmin(np.abs(precision[:-1]-recall[:-1]))].round(4)
plt.axvline(intersection_thr, linestyle='--', color='red')
plt.text(intersection_thr, 0.01, str(intersection_thr), ha='left', color='red')
plt.title("Precision & Recall curve at different thresholds")
plt.xlabel("Threshold values")
plt.ylabel("Precision and Recall values")
plt.legend(loc="upper right")
plt.grid()
plt.show()
```



```
In [315]: threshold_considered = intersection_thr
y_pred_custom = (y_pred_proba>threshold_considered).astype('int')
print(classification_report(y_test,y_pred_custom))
```

	precision	recall	f1-score	support
0	0.92	0.92	0.92	48491
1	0.65	0.65	0.65	11226
accuracy			0.87	59717
macro avg	0.78	0.78	0.78	59717
weighted avg	0.87	0.87	0.87	59717


```
In [316]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred_custom)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



Tradeoff Questions:

- How can we make sure that our model can detect real defaulters and there are less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.

To ensure that our model predicts less false positives, precision of the model should be high.

- Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.

To ensure that our model captures all positives from the dataset or is able to identify all defaulters/non-performing assets, recall should be high. There is a tradeoff between Precision and Recall. As we try to increase precision, recall will be decreased and vice-versa. For this scenario, F1-score should be chosen as a metric to measure the model performance, in a way to achieve the balance b/w precision and recall.

Questions

1. What percentage of customers have fully paid their Loan Amount?
80% of the customers have fully paid their loan and 20% are defaulters
2. Comment about the correlation between Loan Amount and Installment features.
Loan amount and installment are highly correlated(correlation coefficient=0.95) as it is obvious that high loan amount will have high installment amount
3. The majority of people have home ownership as ____.
Most of the people have home ownership as mortgage
4. People with grades 'A' are more likely to fully pay their loan.
True. People with grade A are more likely to fully pay their loan
5. Name the top 2 afforded job titles.
Teacher and Manager are two afforded job titles
6. Thinking from a bank's perspective, which metric should our primary focus be on.
The best metric to consider is F1 score as we need to give importance to both precision and recall. We don't want to miss potential customers and at the same time we also don't want to give loan to defaulters.
7. How does the gap in precision and recall affect the bank?
In case of low recall, false negatives will be high and there will be high number of defaulters. In case of low precision, false positives will be high and LoanTap will lose potential customers. Recall and Precision are contradictory to each other. Increase in one metric will decrease the other. We have to achieve the balance between precision and recall.
8. Which were the features that heavily affected the outcome?
The features zip_code_29597, zip_code_05113, zip_code_00813, annual_inc, loan_amnt, zip_code_86630, zip_code_11650, zip_code_93700, dti, open_acc affected the model outcome heavily
9. Will the results be affected by geographical location? (Yes/No)
Yes. As we can see that zip code is affecting the model which implies that geographical location will affect the results.

Insights

- 80% of the customers have fully paid their loan and 20% are defaulters
- Loan amount and installment are highly correlated as it is obvious that high loan amount will have high installment amount
- Loan taken for short term, i.e. 3 years are most likely to be fully paid back
- Loans with higher amount are less likely to get paid back
- Loans with higher interest rate are less likely to get paid back
- Most of the people have home ownership as mortgage
- Surprisingly, loans which are not verified are more likely to be paid back
- Highest No of loans are issued in October month
- Loan taken as joint application type are more likely to be paid back
- People with grade A are more likely to fully pay their loan
- Loan taken for wedding are more likely to be paid back
- People from zip code 00813, 05113 fully pay back their loans whereas people from zip code 11650, 86630, 93700 are all defaulters
- The features zip_code_29597, zip_code_05113, zip_code_00813, annual_inc, loan_amnt, zip_code_86630, zip_code_11650, zip_code_93700, dti, open_acc affected the model

outcome heavily

- As per the ROC curve and AU-ROC value of 0.9, the model is able to differentiate well between the defaulters and non-defaulters
- As per the PRC and AU-PRC value of 0.77, the model is able to return accurate results as well as return majority of all positive results (high recall)
- For default threshold value, precision for class 1 is 0.46 and recall for class 1 is 0.81. Recall is high and precision is low for this threshold value.
- For this particular case, we have a tradeoff between precision and recall, we try to find best threshold value where precision is equal to recall value. To increase the precision value, we have to bear the reduction in recall.
- For threshold 0.6382, model has a F1-score of 0.65 for class 1. Precision has been increased from 0.47 to 0.65 and Recall has been decreased from 0.81 to 0.65.

Recommendations:

1. LoanTap may provide more short term loans, i.e. for 3 years, without much risk.
2. LoanTap may provide more joint loans and scrutinize more individual and direct pay application types.
3. Analyze carefully the loan applications of customers with grades D, E, F and G. Do not provide them loans or provide smaller loans to these customers for shorter periods.
4. LoanTap may reduce the loans given for small businesses or analyze their applications thoroughly prior giving loans.
5. Do not provide loans to applicants with zip codes 11650, 86630, 93700.
6. Investigate the verifier as the verified loans are not getting paid back and scrutinize the verification process.
7. LoanTap may offer fancy offers like low interest rates or higher loan amount to applicants from zip codes 00813, 05113 to attract more leads.
8. LoanTap may use this model to predict the chances of default before loan approval.